


```
from google.colab import files
uploaded = files.upload()
```

 Choose Files

Dataset of Diabetes .csv


- Dataset of Diabetes .csv(text/csv) - 49511 bytes, last modified: 3/3/2025 - 100% done

Saving Dataset of Diabetes .csv to Dataset of Diabetes .csv


```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.impute import SimpleImputer
from sklearn.preprocessing import OrdinalEncoder, OneHotEncoder
from sklearn.preprocessing import StandardScaler, MinMaxScaler
from scipy import stats

import pandas as pd

# Replace 'your_file.csv' with the name of the file you just uploaded
df = pd.read_csv('Dataset of Diabetes .csv')
df.head() # Display the first few rows
```




	ID	No_Pation	Gender	AGE	Urea	Cr	HbA1c	Chol	TG	HDL	LDL	VLDL	BMI	CLASS
0	502	17975	F	50	4.7	46	4.9	4.2	0.9	2.4	1.4	0.5	24.0	N
1	735	34221	M	26	4.5	62	4.9	3.7	1.4	1.1	2.1	0.6	23.0	N
2	420	47975	F	50	4.7	46	4.9	4.2	0.9	2.4	1.4	0.5	24.0	N
3	680	87656	F	50	4.7	46	4.9	4.2	0.9	2.4	1.4	0.5	24.0	N
4	504	34223	M	33	7.1	46	4.9	4.9	1.0	0.8	2.0	0.4	21.0	N




Next steps: [Generate code with df](#) [View recommended plots](#) [New interactive sheet](#)

```
df.head(10)
```



	ID	No_Pation	Gender	AGE	Urea	Cr	HbA1c	Chol	TG	HDL	LDL	VLDL	BMI	CLASS
0	502	17975	F	50	4.7	46	4.9	4.2	0.9	2.4	1.4	0.5	24.0	N
1	735	34221	M	26	4.5	62	4.9	3.7	1.4	1.1	2.1	0.6	23.0	N
2	420	47975	F	50	4.7	46	4.9	4.2	0.9	2.4	1.4	0.5	24.0	N
3	680	87656	F	50	4.7	46	4.9	4.2	0.9	2.4	1.4	0.5	24.0	N
4	504	34223	M	33	7.1	46	4.9	4.9	1.0	0.8	2.0	0.4	21.0	N
5	634	34224	F	45	2.3	24	4.0	2.9	1.0	1.0	1.5	0.4	21.0	N
6	721	34225	F	50	2.0	50	4.0	3.6	1.3	0.9	2.1	0.6	24.0	N
7	421	34227	M	48	4.7	47	4.0	2.9	0.8	0.9	1.6	0.4	24.0	N
8	670	34229	M	43	2.6	67	4.0	3.8	0.9	2.4	3.7	1.0	21.0	N
9	759	34230	F	32	3.6	28	4.0	3.8	2.0	2.4	3.8	1.0	24.0	N




Next steps: [Generate code with df](#) [View recommended plots](#) [New interactive sheet](#)


```
df.shape
```

(1000, 14)

```
#df.loc[5, 'Age'] = np.nan
#df.loc[10, 'Salary'] = np.nan
df.head(10)
```




	ID	No_Pation	Gender	AGE	Urea	Cr	HbA1c	Chol	TG	HDL	LDL	VLDL	BMI	CLASS
0	502	17975	F	50	4.7	46	4.9	4.2	0.9	2.4	1.4	0.5	24.0	N
1	735	34221	M	26	4.5	62	4.9	3.7	1.4	1.1	2.1	0.6	23.0	N
2	420	47975	F	50	4.7	46	4.9	4.2	0.9	2.4	1.4	0.5	24.0	N
3	680	87656	F	50	4.7	46	4.9	4.2	0.9	2.4	1.4	0.5	24.0	N
4	504	34223	M	33	7.1	46	4.9	4.9	1.0	0.8	2.0	0.4	21.0	N
5	634	34224	F	45	2.3	24	4.0	2.9	1.0	1.0	1.5	0.4	21.0	N
6	721	34225	F	50	2.0	50	4.0	3.6	1.3	0.9	2.1	0.6	24.0	N
7	421	34227	M	48	4.7	47	4.0	2.9	0.8	0.9	1.6	0.4	24.0	N
8	670	34229	M	43	2.6	67	4.0	3.8	0.9	2.4	3.7	1.0	21.0	N
9	759	34230	F	32	3.6	28	4.0	3.8	2.0	2.4	3.8	1.0	24.0	N



Next steps: [Generate code with df](#) [View recommended plots](#) [New interactive sheet](#)

```
print(df.describe())
```



	ID	No_Pation	AGE	Urea	Cr
count	1000.000000	1.000000e+03	1000.000000	1000.000000	1000.000000
mean	340.500000	2.705514e+05	53.528000	5.124743	68.943000
std	240.397673	3.380758e+06	8.799241	2.935165	59.984747
min	1.000000	1.230000e+02	20.000000	0.500000	6.000000
25%	125.750000	2.406375e+04	51.000000	3.700000	48.000000
50%	300.500000	3.439550e+04	55.000000	4.600000	60.000000
75%	550.250000	4.538425e+04	59.000000	5.700000	73.000000
max	800.000000	7.543566e+07	79.000000	38.900000	800.000000

	HbA1c	Chol	TG	HDL	LDL
count	1000.000000	1000.000000	1000.000000	1000.000000	1000.000000
mean	8.281160	4.862820	2.349610	1.204750	2.609790
std	2.534003	1.301738	1.401176	0.660414	1.115102
min	0.900000	0.000000	0.300000	0.200000	0.300000
25%	6.500000	4.000000	1.500000	0.900000	1.800000
50%	8.000000	4.800000	2.000000	1.100000	2.500000
75%	10.200000	5.600000	2.900000	1.300000	3.300000
max	16.000000	10.300000	13.800000	9.900000	9.900000

	VLDL	BMI
count	1000.000000	1000.000000
mean	1.854700	29.578020
std	3.663599	4.962388
min	0.100000	19.000000
25%	0.700000	26.000000
50%	0.900000	30.000000
75%	1.500000	33.000000
max	35.000000	47.750000


```
#Code to Find Missing Values
```

```
# Check for missing values in each column
```

```
missing_values = df.isnull().sum()
```

```
# Display columns with missing values
```

```
print(missing_values[missing_values > 0])
```



```
Series([], dtype: int64)
```

```
import numpy as np
```

```
# Introduce missing values at specific locations
```





```
df.loc[5, 'AGE'] = np.nan # Set missing value for 'AGE' at row index 5
```

```
df.loc[10, 'BMI'] = np.nan # Set missing value for 'BMI' at row index 10
```

```
# Display the first 10 rows to check the changes
```


```
print(df.head(10))
```

```
df
```

	ID	No_Patien	Gender	AGE	Urea	Cr	HbA1c	Cho1	TG	HDL	LDL	VLDL	\		
0	502	17975	F	50.0	4.7	46	4.9	4.2	0.9	2.4	1.4	0.5			
1	735	34221	M	26.0	4.5	62	4.9	3.7	1.4	1.1	2.1	0.6			
2	420	47975	F	50.0	4.7	46	4.9	4.2	0.9	2.4	1.4	0.5			
3	680	87656	F	50.0	4.7	46	4.9	4.2	0.9	2.4	1.4	0.5			
4	504	34223	M	33.0	7.1	46	4.9	4.9	1.0	0.8	2.0	0.4			
5	634	34224	F	NaN	2.3	24	4.0	2.9	1.0	1.0	1.5	0.4			
6	721	34225	F	50.0	2.0	50	4.0	3.6	1.3	0.9	2.1	0.6			
7	421	34227	M	48.0	4.7	47	4.0	2.9	0.8	0.9	1.6	0.4			
8	670	34229	M	43.0	2.6	67	4.0	3.8	0.9	2.4	3.7	1.0			
9	759	34230	F	32.0	3.6	28	4.0	3.8	2.0	2.4	3.8	1.0			
BMI CLASS															
0	24.0	N													
1	23.0	N													
2	24.0	N													
3	24.0	N													
4	21.0	N													
5	21.0	N													
6	24.0	N													
7	24.0	N													
8	21.0	N													
9	24.0	N													
	ID	No_Patien	Gender	AGE	Urea	Cr	HbA1c	Cho1	TG	HDL	LDL	VLDL	BMI	CLASS	
0	502	17975	F	50.0	4.7	46	4.9	4.2	0.9	2.4	1.4	0.5	24.0	N	
1	735	34221	M	26.0	4.5	62	4.9	3.7	1.4	1.1	2.1	0.6	23.0	N	
2	420	47975	F	50.0	4.7	46	4.9	4.2	0.9	2.4	1.4	0.5	24.0	N	
3	680	87656	F	50.0	4.7	46	4.9	4.2	0.9	2.4	1.4	0.5	24.0	N	
4	504	34223	M	33.0	7.1	46	4.9	4.9	1.0	0.8	2.0	0.4	21.0	N	
...	
995	200	454317	M	71.0	11.0	97	7.0	7.5	1.7	1.2	1.8	0.6	30.0	Y	
996	671	876534	M	31.0	3.0	60	12.3	4.1	2.2	0.7	2.4	15.4	37.2	Y	
997	669	87654	M	30.0	7.1	81	6.7	4.1	1.1	1.2	2.4	8.1	27.4	Y	
998	99	24004	M	38.0	5.8	59	6.7	5.3	2.0	1.6	2.9	14.0	40.5	Y	
999	248	24054	M	54.0	5.0	67	6.9	3.8	1.7	1.1	3.0	0.7	33.0	Y	
1000 rows × 14 columns															

Next steps: [Generate code with df](#) [View recommended plots](#) [New interactive sheet](#)

```
print(df.describe())
```

		ID	No_Patien	AGE	Urea	Cr	\
count	1000.000000	1.000000e+03	999.000000	1000.000000	1000.000000	1000.000000	
mean	340.500000	2.705514e+05	53.536537	5.124743	68.943000		
std	240.397673	3.380758e+06	8.799504	2.935165	59.984747		
min	1.000000	1.230000e+02	20.000000	0.500000	6.000000		
25%	125.750000	2.406375e+04	51.000000	3.700000	48.000000		
50%	300.500000	3.439550e+04	55.000000	4.600000	60.000000		
75%	550.250000	4.538425e+04	59.000000	5.700000	73.000000		
max	800.000000	7.543566e+07	79.000000	38.900000	800.000000		
	HbA1c	Cho1	TG	HDL	LDL	\	
count	1000.000000	1000.000000	1000.000000	1000.000000	1000.000000		
mean	8.281160	4.862820	2.349610	1.204750	2.609790		
std	2.534003	1.301738	1.401176	0.660414	1.115102		
min	0.900000	0.000000	0.300000	0.200000	0.300000		
25%	6.500000	4.000000	1.500000	0.900000	1.800000		
50%	8.000000	4.800000	2.000000	1.100000	2.500000		
75%	10.200000	5.600000	2.900000	1.300000	3.300000		
max	16.000000	10.300000	13.800000	9.900000	9.900000		
	VLDL	BMI					
count	1000.000000	999.000000					
mean	1.854700	29.584605					
std	3.663599	4.960501					
min	0.100000	19.000000					
25%	0.700000	26.000000					
50%	0.900000	30.000000					
75%	1.500000	33.000000					
max	35.000000	47.750000					

```

#Code to Find Missing Values
# Check for missing values in each column
missing_values = df.isnull().sum()

# Display columns with missing values
print(missing_values[missing_values > 0])

AGE      1
BMI      1
dtype: int64

#Set the values to some value (zero, the mean, the median, etc.).
# Step 1: Create an instance of SimpleImputer with the median strategy for Age and mean strategy for Salary
imputer1 = SimpleImputer(strategy="median")
imputer2 = SimpleImputer(strategy="mean")

df_copy=df

# Step 2: Fit the imputer on the "Age" and "Salary" column
# Note: SimpleImputer expects a 2D array, so we reshape the column
imputer1.fit(df_copy[["AGE"]])
imputer2.fit(df_copy[["BMI"]])

# Step 3: Transform (fill) the missing values in the "Age" and "Salary" column
df_copy["AGE"] = imputer1.transform(df[["AGE"]])
df_copy["BMI"] = imputer2.transform(df[["BMI"]])

# Verify that there are no missing values left
print(df_copy["AGE"].isnull().sum())
print(df_copy["BMI"].isnull().sum())

0
0

import pandas as pd
from sklearn.preprocessing import OrdinalEncoder, OneHotEncoder

# Normalize the Gender column to be consistent (uppercase in this case)
df['Gender'] = df['Gender'].str.upper() # Convert to uppercase
df['CLASS'] = df['CLASS'].str.upper()
# Initialize OrdinalEncoder for 'Gender' column
ordinal_encoder = OrdinalEncoder(categories=[["F", "M"]]) # Encoding 'F' as 0, 'M' as 1

# Fit and transform the data in the 'Gender' column
df["Gender_Encoded"] = ordinal_encoder.fit_transform(df[["Gender"]])

# Initialize OneHotEncoder for the 'City' column (if the column exists)
# You should replace "City" with the actual column name in your dataset
onehot_encoder = OneHotEncoder()

# Fit and transform the "City" column (replace 'City' with the actual name of the column)
if 'CLASS' in df.columns:
    encoded_data = onehot_encoder.fit_transform(df[["CLASS"]])

    # Convert the sparse matrix to a dense array
    encoded_array = encoded_data.toarray()

    # Convert to DataFrame for better visualization
    encoded_df = pd.DataFrame(encoded_array, columns=onehot_encoder.get_feature_names_out(["CLASS"]))

    # Concatenate the one-hot encoded columns with the original DataFrame
    df_encoded = pd.concat([df, encoded_df], axis=1)

    # Drop the original 'City' column as it is now encoded
    df_encoded.drop("CLASS", axis=1, inplace=True)

# If there is no 'City' column, proceed with just encoding 'Gender'
else:
    df_encoded = df.copy()

# Drop the original 'Gender' column
df_encoded.drop("Gender", axis=1, inplace=True)

# Display the first few rows of the encoded dataframe
print(df_encoded.head())
df_encoded

```

1000 rows x 18 columns

```
import pandas as pd
from sklearn.preprocessing import MinMaxScaler

# Initialize the MinMaxScaler
normalizer = MinMaxScaler()

# Apply the MinMaxScaler to the 'AGE' column only
df_encoded[['AGE']] = normalizer.fit_transform(df_encoded[['AGE']])

# Display the first few rows to verify the transformation
print(df_encoded.head())
df_encoded
```

	ID	No_Pation	AGE	Urea	Cr	HbA1c	Chol	TG	HDL	LDL	VLDL	BMI	\					
0	502	17975	0.508475	4.7	46	4.9	4.2	0.9	2.4	1.4	0.5	24.0						
1	735	34221	0.101695	4.5	62	4.9	3.7	1.4	1.1	2.1	0.6	23.0						
2	420	47975	0.508475	4.7	46	4.9	4.2	0.9	2.4	1.4	0.5	24.0						
3	680	87656	0.508475	4.7	46	4.9	4.2	0.9	2.4	1.4	0.5	24.0						
4	504	34223	0.220339	7.1	46	4.9	4.9	1.0	0.8	2.0	0.4	21.0						
	Gender_Encoded	CLASS_N	CLASS_N	CLASS_P	CLASS_Y	CLASS_Y												
0	0.0	1.0	0.0	0.0	0.0	0.0												
1	1.0	1.0	0.0	0.0	0.0	0.0												
2	0.0	1.0	0.0	0.0	0.0	0.0												
3	0.0	1.0	0.0	0.0	0.0	0.0												
4	1.0	1.0	0.0	0.0	0.0	0.0												
	ID	No_Pation	AGE	Urea	Cr	HbA1c	Chol	TG	HDL	LDL	VLDL	BMI	Gender_Encoded	CLASS_N	CLASS_N	CLASS_P	CLASS_Y	CLASS_Y
0	502	17975	0.508475	4.7	46	4.9	4.2	0.9	2.4	1.4	0.5	24.0	0.0	1.0	0.0	0.0	0.0	0.0
1	735	34221	0.101695	4.5	62	4.9	3.7	1.4	1.1	2.1	0.6	23.0	1.0	1.0	0.0	0.0	0.0	0.0
2	420	47975	0.508475	4.7	46	4.9	4.2	0.9	2.4	1.4	0.5	24.0	0.0	1.0	0.0	0.0	0.0	0.0
3	680	87656	0.508475	4.7	46	4.9	4.2	0.9	2.4	1.4	0.5	24.0	0.0	1.0	0.0	0.0	0.0	0.0
4	504	34223	0.220339	7.1	46	4.9	4.9	1.0	0.8	2.0	0.4	21.0	1.0	1.0	0.0	0.0	0.0	0.0
...
995	200	454317	0.864407	11.0	97	7.0	7.5	1.7	1.2	1.8	0.6	30.0	1.0	0.0	0.0	0.0	1.0	0.0
996	671	876534	0.186441	3.0	60	12.3	4.1	2.2	0.7	2.4	15.4	37.2	1.0	0.0	0.0	0.0	0.0	1.0
997	669	87654	0.169492	7.1	81	6.7	4.1	1.1	1.2	2.4	8.1	27.4	1.0	0.0	0.0	0.0	0.0	1.0
998	99	24004	0.305085	5.8	59	6.7	5.3	2.0	1.6	2.9	14.0	40.5	1.0	0.0	0.0	0.0	0.0	1.0
999	248	24054	0.576271	5.0	67	6.9	3.8	1.7	1.1	3.0	0.7	33.0	1.0	0.0	0.0	0.0	0.0	1.0

1000 rows × 18 columns

Next steps: [Generate code with df_encoded](#) [View recommended plots](#) [New interactive sheet](#)

```
#Data Transformation
# Min-Max Scaler/Normalization (range 0-1)
#Pros: Keeps all data between 0 and 1; ideal for distance-based models.
#Cons: Can distort data distribution, especially with extreme outliers.
normalizer = MinMaxScaler()
df_encoded[['TG']] = normalizer.fit_transform(df_encoded[['TG']])
df_encoded.head()
```

	ID	No_Pation	AGE	Urea	Cr	HbA1c	Chol	TG	HDL	LDL	VLDL	BMI	Gender_Encoded	CLASS_N	CLASS_N	CLASS_P	CLASS_Y	CLASS_Y
0	502	17975	50.0	4.7	46	4.9	4.2	0.044444	2.4	1.4	0.5	24.0	0.0	1.0	0.0	0.0	0.0	0.0
1	735	34221	26.0	4.5	62	4.9	3.7	0.081481	1.1	2.1	0.6	23.0	1.0	1.0	0.0	0.0	0.0	0.0
2	420	47975	50.0	4.7	46	4.9	4.2	0.044444	2.4	1.4	0.5	24.0	0.0	1.0	0.0	0.0	0.0	0.0
3	680	87656	50.0	4.7	46	4.9	4.2	0.044444	2.4	1.4	0.5	24.0	0.0	1.0	0.0	0.0	0.0	0.0
4	504	34223	33.0	7.1	46	4.9	4.9	0.051852	0.8	2.0	0.4	21.0	1.0	1.0	0.0	0.0	0.0	0.0

Next steps: [Generate code with df_encoded](#) [View recommended plots](#) [New interactive sheet](#)

```
# Standardization (mean=0, variance=1)
#Pros: Works well for normally distributed data; suitable for many models.
#Cons: Sensitive to outliers.
scaler = StandardScaler()
df_encoded[['AGE']] = scaler.fit_transform(df_encoded[['AGE']])
df_encoded.head()
```

	ID	No_Pation	AGE	Urea	Cr	HbA1c	Chol	TG	HDL	LDL	VLDL	BMI	Gender_Encoded	CLASS_N	CLASS_N	CLASS_P	CLASS_Y	CL
0	502	17975	-0.402465	4.7	46	4.9	4.2	0.044444	2.4	1.4	0.5	24.0	0.0	1.0	0.0	0.0	0.0	
1	735	34221	-3.132585	4.5	62	4.9	3.7	0.081481	1.1	2.1	0.6	23.0	1.0	1.0	0.0	0.0	0.0	
2	420	47975	-0.402465	4.7	46	4.9	4.2	0.044444	2.4	1.4	0.5	24.0	0.0	1.0	0.0	0.0	0.0	
3	680	87656	-0.402465	4.7	46	4.9	4.2	0.044444	2.4	1.4	0.5	24.0	0.0	1.0	0.0	0.0	0.0	
4	504	34223	-2.336300	7.1	46	4.9	4.9	0.051852	0.8	2.0	0.4	21.0	1.0	1.0	0.0	0.0	0.0	

Next steps:

[Generate code with df_encoded](#)[View recommended plots](#)[New interactive sheet](#)

```
#Removing Outliers
# Outlier Detection and Treatment using IQR
#Pros: Simple and effective for mild outliers.
#Cons: May overly reduce variation if there are many extreme outliers.
df_encoded_copy1=df_encoded
df_encoded_copy2=df_encoded
df_encoded_copy3=df_encoded

Q1 = df_encoded_copy1['Chol'].quantile(0.25)
Q3 = df_encoded_copy1['Chol'].quantile(0.75)
IQR = Q3 - Q1
lower_bound = Q1 - 1.5 * IQR
upper_bound = Q3 + 1.5 * IQR
df_encoded_copy1['Chol'] = np.where(df_encoded_copy1['Chol'] > upper_bound, upper_bound,
np.where(df_encoded_copy1['Chol'] < lower_bound, lower_bound, df_encoded_copy1['Chol']))

print(df_encoded_copy1.head())
```

```
↗
  ID  No_Patien    AGE  Urea  Cr  HbA1c  Chol    TG  HDL  LDL  VLDL  \
0  502    17975 -0.402465  4.7  46    4.9    4.2  0.044444  2.4  1.4  0.5
1  735    34221 -3.132585  4.5  62    4.9    3.7  0.081481  1.1  2.1  0.6
2  420    47975 -0.402465  4.7  46    4.9    4.2  0.044444  2.4  1.4  0.5
3  680    87656 -0.402465  4.7  46    4.9    4.2  0.044444  2.4  1.4  0.5
4  504    34223 -2.336300  7.1  46    4.9    4.9  0.051852  0.8  2.0  0.4
```

```
      BMI  Gender_Encoded  CLASS_N  CLASS_N  CLASS_P  CLASS_Y  CLASS_Y
0  24.0          0.0        1.0        0.0        0.0        0.0        0.0
1  23.0          1.0        1.0        0.0        0.0        0.0        0.0
2  24.0          0.0        1.0        0.0        0.0        0.0        0.0
3  24.0          0.0        1.0        0.0        0.0        0.0        0.0
4  21.0          1.0        1.0        0.0        0.0        0.0        0.0
```

```
#Removing Outliers
# Z-score method
#Pros: Good for normally distributed data.
#Cons: Not suitable for non-normal data; may miss outliers in skewed distributions.
```

```
df_encoded_copy2['Chol_zscore'] = stats.zscore(df_encoded_copy2['Chol'])
df_encoded_copy2['Chol'] = np.where(df_encoded_copy2['Chol_zscore'].abs() > 3, np.nan, df_encoded_copy2['Chol']) # Replace outliers with NaN
print(df_encoded_copy2.head())
```

```
↗
  ID  No_Patien    AGE  Urea  Cr  HbA1c  Chol    TG  HDL  LDL  VLDL  \
0  502    17975 -0.402465  4.7  46    4.9    4.2  0.044444  2.4  1.4  0.5
1  735    34221 -3.132585  4.5  62    4.9    3.7  0.081481  1.1  2.1  0.6
2  420    47975 -0.402465  4.7  46    4.9    4.2  0.044444  2.4  1.4  0.5
3  680    87656 -0.402465  4.7  46    4.9    4.2  0.044444  2.4  1.4  0.5
4  504    34223 -2.336300  7.1  46    4.9    4.9  0.051852  0.8  2.0  0.4
```

```
      BMI  Gender_Encoded  CLASS_N  CLASS_N  CLASS_P  CLASS_Y  CLASS_Y  \
0  24.0          0.0        1.0        0.0        0.0        0.0        0.0
1  23.0          1.0        1.0        0.0        0.0        0.0        0.0
2  24.0          0.0        1.0        0.0        0.0        0.0        0.0
3  24.0          0.0        1.0        0.0        0.0        0.0        0.0
4  21.0          1.0        1.0        0.0        0.0        0.0        0.0
```

```
      Chol_zscore
0   -0.532005
1   -0.945425
2   -0.532005
3   -0.532005
4    0.046783
```

```
#Removing Outliers
# Median replacement for outliers
#Pros: Keeps distribution shape intact, useful when capping isn't feasible.
#Cons: May distort data if outliers represent real phenomena.
df_encoded_copy3['Chol_zscore'] = stats.zscore(df_encoded_copy3['Chol'])
median_salary = df_encoded_copy3['Chol'].median()
df_encoded_copy3['Chol'] = np.where(df_encoded_copy3['Chol'].abs() > 3, median_salary, df_encoded_copy3['Chol'])
print(df_encoded_copy3.head())
```

```
↗
  ID  No_Patien    AGE  Urea  Cr  HbA1c  Chol    TG  HDL  LDL  VLDL  \
0  502    17975 -0.402465  4.7  46    4.9    4.8  0.044444  2.4  1.4  0.5
1  735    34221 -3.132585  4.5  62    4.9    4.8  0.081481  1.1  2.1  0.6
```

2 420 47975 -0.402465 4.7 46 4.9 4.8 0.044444 2.4 1.4 0.5