


```
from google.colab import files
uploaded = files.upload()

 Choose Files adult.csv


- adult.csv(text/csv) - 5326368 bytes, last modified: 3/3/2025 - 100% done




Saving adult.csv to adult.csv



import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.impute import SimpleImputer
from sklearn.preprocessing import OrdinalEncoder, OneHotEncoder
from sklearn.preprocessing import StandardScaler, MinMaxScaler
from scipy import stats

import pandas as pd


# Replace 'your_file.csv' with the name of the file you just uploaded
df = pd.read_csv('adult.csv')
df.head() # Display the first few rows
```



	age	workclass	fnlwtg	education	educational-num	marital-status	occupation	relationship	race	gender	capital-gain	capital-loss	hours-per-week	native-country
0	25	Private	226802	11th	7	Never-married	Machine-op-inspct	Own-child	Black	Male	0	0	40	United States
1	38	Private	89814	HS-grad	9	Married-civ-spouse	Farming-fishing	Husband	White	Male	0	0	50	United States

Next steps: [Generate code with df](#) [View recommended plots](#) [New interactive sheet](#)


```
df.head(10)
```



	age	workclass	fnlwtg	education	educational-num	marital-status	occupation	relationship	race	gender	capital-gain	capital-loss	hours-per-week	native-country
0	25	Private	226802	11th	7	Never-married	Machine-op-inspct	Own-child	Black	Male	0	0	40	United States
1	38	Private	89814	HS-grad	9	Married-civ-spouse	Farming-fishing	Husband	White	Male	0	0	50	United States
2	28	Local-gov	336951	Assoc-acdm	12	Married-civ-spouse	Protective-serv	Husband	White	Male	0	0	40	United States
3	44	Private	160323	Some-college	10	Married-civ-spouse	Machine-op-inspct	Husband	Black	Male	7688	0	40	United States
4	18	2	103407	Some-	10	Never-	2	Own-child	White	Female	0	0	20	United States


Next steps: [Generate code with df](#) [View recommended plots](#) [New interactive sheet](#)

```
df.shape
```



(48842, 15)

```
print(df.describe())
```



	age	fnlwtg	educational-num	capital-gain \
count	48842.000000	4.884200e+04	48842.000000	48842.000000
mean	38.643585	1.896641e+05	10.078089	1079.067626
std	13.710510	1.056040e+05	2.570973	7452.019058
min	17.000000	1.228500e+04	1.000000	0.000000
25%	28.000000	1.175505e+05	9.000000	0.000000
50%	37.000000	1.781445e+05	10.000000	0.000000

75%	48.000000	2.376420e+05	12.000000	0.000000
max	90.000000	1.490400e+06	16.000000	99999.000000

	capital-loss	hours-per-week
count	48842.000000	48842.000000
mean	87.502314	40.422382
std	403.004552	12.391444
min	0.000000	1.000000
25%	0.000000	40.000000
50%	0.000000	40.000000
75%	0.000000	45.000000
max	4356.000000	99.000000

```
#Code to Find Missing Values
```

```
# Check for missing values in each column
```

```
missing_values = df.isnull().sum()
```

```
# Display columns with missing values
```

```
print(missing_values[missing_values > 0])
```

```
Series([], dtype: int64)
```

```
import numpy as np
```

```
# Introduce missing values at specific locations
```

```
df.loc[5, 'educational-num'] = np.nan # Set missing value for 'AGE' at row index 5
```

```
df.loc[7, 'age'] = np.nan # Set missing value for 'BMI' at row index 10
```

```
# Display the first 10 rows to check the changes
```

```
print(df.head(10))
```

```
df
```



0	25.0	Private	226802	11th	7.0
1	38.0	Private	89814	HS-grad	9.0
2	28.0	Local-gov	336951	Assoc-acdm	12.0
3	44.0	Private	160323	Some-college	10.0
4	18.0	?	103497	Some-college	10.0
5	34.0	Private	198693	10th	NaN
6	29.0	?	227026	HS-grad	9.0
7	NaN	Self-emp-not-inc	104626	Prof-school	15.0
8	24.0	Private	369667	Some-college	10.0
9	55.0	Private	104996	7th-8th	4.0

	marital-status	occupation	relationship	race	gender
0	Never-married	Machine-op-inspct	Own-child	Black	Male
1	Married-civ-spouse	Farming-fishing	Husband	White	Male
2	Married-civ-spouse	Protective-serv	Husband	White	Male
3	Married-civ-spouse	Machine-op-inspct	Husband	Black	Male
4	Never-married	?	Own-child	White	Female
5	Never-married	Other-service	Not-in-family	White	Male
6	Never-married	?	Unmarried	Black	Male
7	Married-civ-spouse	Prof-specialty	Husband	White	Male
8	Never-married	Other-service	Unmarried	White	Female
9	Married-civ-spouse	Craft-repair	Husband	White	Male

	capital-gain	capital-loss	hours-per-week	native-country	income
0	0	0	40	United-States	<=50K
1	0	0	50	United-States	<=50K
2	0	0	40	United-States	>50K
3	7688	0	40	United-States	>50K
4	0	0	30	United-States	<=50K
5	0	0	30	United-States	<=50K
6	0	0	40	United-States	<=50K
7	3103	0	32	United-States	>50K
8	0	0	40	United-States	<=50K
9	0	0	10	United-States	<=50K

	age	workclass	fnlwgt	education	educational-num	marital-status	occupation	relationship	race	gender	capital-gain	capital-loss	hours-per-week
0	25.0	Private	226802	11th	7.0	Never-married	Machine-op-inspct	Own-child	Black	Male	0	0	40
1	38.0	Private	89814	HS-grad	9.0	Married-civ-spouse	Farming-fishing	Husband	White	Male	0	0	50
2	28.0	Local-gov	336951	Assoc-acdm	12.0	Married-civ-spouse	Protective-serv	Husband	White	Male	0	0	40
3	44.0	Private	160323	Some-college	10.0	Married-civ-spouse	Machine-op-inspct	Husband	Black	Male	7688	0	40
4	18.0	?	103497	Some-college	10.0	Never-married	?	Own-child	White	Female	0	0	30
...
48837	27.0	Private	257302	Assoc-acdm	12.0	Married-civ-spouse	Tech-support	Wife	White	Female	0	0	38

Next steps:

[Generate code with df](#)
[View recommended plots](#)
[New interactive sheet](#)

print(df.describe())



	age	fnlwgt	educational-num	capital-gain
count	48841.000000	4.884200e+04	48841.000000	48842.000000
mean	38.643087	1.896641e+05	10.078172	1079.067626
std	13.710207	1.056040e+05	2.570933	7452.019058
min	17.000000	1.228500e+04	1.000000	0.000000
25%	28.000000	1.175505e+05	9.000000	0.000000
50%	37.000000	1.781445e+05	10.000000	0.000000
75%	48.000000	2.376420e+05	12.000000	0.000000
max	90.000000	1.490400e+06	16.000000	99999.000000

	capital-loss	hours-per-week
count	48842.000000	48842.000000
mean	87.502314	40.422382
std	403.004552	12.391444
min	0.000000	1.000000
25%	0.000000	40.000000
50%	0.000000	40.000000
75%	0.000000	45.000000
max	4356.000000	99.000000

```

#Code to Find Missing Values
# Check for missing values in each column
missing_values = df.isnull().sum()

# Display columns with missing values
print(missing_values[missing_values > 0])

↩ age          1
  educational-num  1
  dtype: int64

#Set the values to some value (zero, the mean, the median, etc.).
# Step 1: Create an instance of SimpleImputer with the median strategy for Age and mean strategy for Salary
imputer1 = SimpleImputer(strategy="median")
imputer2 = SimpleImputer(strategy="mean")

df_copy=df

# Step 2: Fit the imputer on the "Age" and "Salary"column
# Note: SimpleImputer expects a 2D array, so we reshape the column
imputer1.fit(df_copy[["educational-num"]])
imputer2.fit(df_copy[["age"]])

# Step 3: Transform (fill) the missing values in the "Age" and "Salary" column
df_copy["educational-num"] = imputer1.transform(df[["educational-num"]])
df_copy["age"] = imputer2.transform(df[["age"]])

# Verify that there are no missing values left
print(df_copy["educational-num"].isnull().sum())
print(df_copy["age"].isnull().sum())

↩ 0
  0

import pandas as pd
from sklearn.preprocessing import OrdinalEncoder, OneHotEncoder

# Normalize the Gender column to be consistent (uppercase in this case)
df['gender'] = df['gender'].str.upper() # Convert to uppercase
df['race'] = df['race'].str.upper()
# Initialize OrdinalEncoder for 'Gender' column
ordinal_encoder = OrdinalEncoder(categories=[["FEMALE", "MALE"]]) # Encoding 'F' as 0, 'M' as 1

# Fit and transform the data in the 'Gender' column
df["gender_encoded"] = ordinal_encoder.fit_transform(df[["gender"]])

# Initialize OneHotEncoder for the 'City' column (if the column exists)
# You should replace "City" with the actual column name in your dataset
onehot_encoder = OneHotEncoder()

# Fit and transform the "City" column (replace 'City' with the actual name of the column)
if 'race' in df.columns:
    encoded_data = onehot_encoder.fit_transform(df[["race"]])

    # Convert the sparse matrix to a dense array
    encoded_array = encoded_data.toarray()

    # Convert to DataFrame for better visualization
    encoded_df = pd.DataFrame(encoded_array, columns=onehot_encoder.get_feature_names_out(["race"]))

    # Concatenate the one-hot encoded columns with the original DataFrame
    df_encoded = pd.concat([df, encoded_df], axis=1)

    # Drop the original 'City' column as it is now encoded
    df_encoded.drop("race", axis=1, inplace=True)

# If there is no 'City' column, proceed with just encoding 'Gender'
else:
    df_encoded = df.copy()

# Drop the original 'Gender' column
df_encoded.drop("gender", axis=1, inplace=True)

# Display the first few rows of the encoded dataframe

```

```
print(df_encoded.head())
df_encoded
```

	age	workclass	fnlwgt	education	educational-num	marital-status	\															
0	25.0	Private	226802	11th	7.0	Never-married																
1	38.0	Private	89814	HS-grad	9.0	Married-civ-spouse																
2	28.0	Local-gov	336951	Assoc-acdm	12.0	Married-civ-spouse																
3	44.0	Private	160323	Some-college	10.0	Married-civ-spouse																
4	18.0	?	103497	Some-college	10.0	Never-married																
	occupation	relationship	capital-gain	capital-loss	hours-per-week	\																
0	Machine-op-inspct	Own-child	0	0	40																	
1	Farming-fishing	Husband	0	0	50																	
2	Protective-serv	Husband	0	0	40																	
3	Machine-op-inspct	Husband	7688	0	40																	
4	?	Own-child	0	0	30																	
	native-country	income	gender_Encoded	race_AMER-INDIAN-ESKIMO	\																	
0	United-States	<=50K	1.0	0.0																		
1	United-States	<=50K	1.0	0.0																		
2	United-States	>50K	1.0	0.0																		
3	United-States	>50K	1.0	0.0																		
4	United-States	<=50K	0.0	0.0																		
	race_ASIAN-PAC-ISLANDER	race_BLACK	race_OTHER	race_WHITE																		
0	0.0	1.0	0.0	0.0																		
1	0.0	0.0	0.0	1.0																		
2	0.0	0.0	0.0	1.0																		
3	0.0	1.0	0.0	0.0																		
4	0.0	0.0	0.0	1.0																		
	age	workclass	fnlwgt	education	educational-num	marital-status	occupation	relationship	capital-gain	capital-loss	hours-per-week	native-country	income									
0	25.0	Private	226802	11th	7.0	Never-married	Machine-op-inspct	Own-child	0	0	40	United-States	<=50K									
1	38.0	Private	89814	HS-grad	9.0	Married-civ-spouse	Farming-fishing	Husband	0	0	50	United-States	<=50K									
2	28.0	Local-gov	336951	Assoc-acdm	12.0	Married-civ-spouse	Protective-serv	Husband	0	0	40	United-States	>50K									
3	44.0	Private	160323	Some-college	10.0	Married-civ-spouse	Machine-op-inspct	Husband	7688	0	40	United-States	>50K									
4	18.0	?	103497	Some-college	10.0	Never-married	?	Own-child	0	0	30	United-States	<=50K									
...									
48837	27.0	Private	257302	Assoc-acdm	12.0	Married-civ-spouse	Tech-support	Wife	0	0	38	United-States	<=50K									
48838	40.0	Private	154374	HS-grad	9.0	Married-civ-spouse	Machine-op-inspct	Husband	0	0	40	United-States	>50K									
48839	58.0	Private	151910	HS-grad	9.0	Widowed	Adm-clerical	Unmarried	0	0	40	United-States	<=50K									
48840	22.0	Private	201490	HS-grad	9.0	Never-married	Adm-clerical	Own-child	0	0	20	United-States	<=50K									

Next steps: [Generate code with df_encoded](#) [View recommended plots](#) [New interactive sheet](#)

```
import pandas as pd
from sklearn.preprocessing import MinMaxScaler

# Initialize the MinMaxScaler
normalizer = MinMaxScaler()

# Apply the MinMaxScaler to the 'AGE' column only
df_encoded[['hours-per-week']] = normalizer.fit_transform(df_encoded[['hours-per-week']])

# Display the first few rows to verify the transformation
print(df_encoded.head())
df_encoded
```



```
4 18.0      ? 103497 Some-college      10.0      never-married
```

```

      occupation relationship capital-gain capital-loss hours-per-week \
0  Machine-op-inspct Own-child      0      0      0.397959
1  Farming-fishing   Husband      0      0      0.500000
2  Protective-serv   Husband      0      0      0.397959
3  Machine-op-inspct Husband    7688      0      0.397959
4      ? Own-child      0      0      0.295918
```

```

native-country income gender_Encoded race_AMER-INDIAN-ESKIMO \
0 United-States <=50K      1.0      0.0
1 United-States <=50K      1.0      0.0
2 United-States >50K      1.0      0.0
3 United-States >50K      1.0      0.0
4 United-States <=50K      0.0      0.0
```

```

race_ASIAN-PAC-ISLANDER race_BLACK race_OTHER race_WHITE
0      0.0      1.0      0.0      0.0
1      0.0      0.0      0.0      1.0
2      0.0      0.0      0.0      1.0
3      0.0      1.0      0.0      0.0
4      0.0      0.0      0.0      1.0
```

	age	workclass	fnlwgt	education	educational-num	marital-status	occupation	relationship	capital-gain	capital-loss	hours-per-week	native-country	income
0	25.0	Private	226802	11th	7.0	Never-married	Machine-op-inspct	Own-child	0	0	0.397959	United-States	<=50K
1	38.0	Private	89814	HS-grad	9.0	Married-civ-spouse	Farming-fishing	Husband	0	0	0.500000	United-States	<=50K
2	28.0	Local-gov	336951	Assoc-acdm	12.0	Married-civ-spouse	Protective-serv	Husband	0	0	0.397959	United-States	>50K
3	44.0	Private	160323	Some-college	10.0	Married-civ-spouse	Machine-op-inspct	Husband	7688	0	0.397959	United-States	>50K
4	18.0	?	103497	Some-college	10.0	Never-married	?	Own-child	0	0	0.295918	United-States	<=50K
...
48837	27.0	Private	257302	Assoc-acdm	12.0	Married-civ-spouse	Tech-support	Wife	0	0	0.377551	United-States	<=50K
48838	40.0	Private	154374	HS-grad	9.0	Married-civ-spouse	Machine-op-inspct	Husband	0	0	0.397959	United-States	>50K
48839	58.0	Private	151910	HS-grad	9.0	Widowed	Adm-clerical	Unmarried	0	0	0.397959	United-States	<=50K
48840	22.0	Private	201490	HS-grad	9.0	Never-married	Adm-clerical	Own-child	0	0	0.193878	United-States	<=50K
48841	52.0	Private	287927	HS-grad	9.0	Married-civ-spouse	Managerial	Wife	15024	0	0.397959	United-States	>50K

```

# Standardization (mean=0, variance=1)
#Pros: Works well for normally distributed data; suitable for many models.
#Cons: Sensitive to outliers.
scaler = StandardScaler()
df_encoded[['age']] = scaler.fit_transform(df_encoded[['age']])
df_encoded.head()
```



	age	workclass	fnlwgt	education	educational-num	marital-status	occupation	relationship	capital-gain	capital-loss	hours-per-week	native-country	income
0	-0.995125	Private	226802	11th	7.0	Never-married	Machine-op-inspct	Own-child	0	0	0.397959	United-States	<=50K
1	-0.046907	Private	89814	HS-grad	9.0	Married-civ-spouse	Farming-fishing	Husband	0	0	0.500000	United-States	<=50K
2	-0.776305	Local-gov	336951	Assoc-acdm	12.0	Married-civ-spouse	Protective-serv	Husband	0	0	0.397959	United-States	>50K
3	0.390732	Private	160323	Some-college	10.0	Married-civ-spouse	Machine-op-inspct	Husband	7688	0	0.397959	United-States	>50K

Next steps:

[Generate code with df_encoded](#)[View recommended plots](#)[New interactive sheet](#)

#Removing Outliers

Outlier Detection and Treatment using IQR

#Pros: Simple and effective for mild outliers.

#Cons: May overly reduce variation if there are many extreme outliers.

df_encoded_copy1=df_encoded

df_encoded_copy2=df_encoded

df_encoded_copy3=df_encoded

Q1 = df_encoded_copy1['fnlwgt'].quantile(0.25)

Q3 = df_encoded_copy1['fnlwgt'].quantile(0.75)

IQR = Q3 - Q1

lower_bound = Q1 - 1.5 * IQR

upper_bound = Q3 + 1.5 * IQR

```
df_encoded_copy1['fnlwgt'] = np.where(df_encoded_copy1['fnlwgt'] > upper_bound, upper_bound,
                                     np.where(df_encoded_copy1['fnlwgt'] < lower_bound, lower_bound, df_encoded_copy1['fnlwgt']))
```

print(df_encoded_copy1.head())



	age	workclass	fnlwgt	education	educational-num	marital-status	occupation	relationship	capital-gain	capital-loss	hours-per-week	native-country	income	gender_Encoded	race_AMER-INDIAN-ESKIMO	race_ASIAN-PAC-ISLANDER	race_BLACK	race_OTHER	race_WHITE
0	-0.995125	Private	226802.0	11th	7.0	Never-married	Machine-op-inspct	Own-child	0	0	0.397959	United-States	<=50K	1.0	0.0	0.0	1.0	0.0	0.0
1	-0.046907	Private	89814.0	HS-grad	9.0	Married-civ-spouse	Farming-fishing	Husband	0	0	0.500000	United-States	<=50K	1.0	0.0	0.0	0.0	0.0	1.0
2	-0.776305	Local-gov	336951.0	Assoc-acdm	12.0	Married-civ-spouse	Protective-serv	Husband	0	0	0.397959	United-States	>50K	1.0	0.0	0.0	0.0	0.0	1.0
3	0.390732	Private	160323.0	Some-college	10.0	Married-civ-spouse	Machine-op-inspct	Husband	7688	0	0.397959	United-States	>50K	1.0	0.0	0.0	1.0	0.0	0.0
4	-1.505704	?	103497.0	Some-college	10.0	Never-married	?	Own-child	0	0	0.295918	United-States	<=50K	0.0	0.0	0.0	0.0	0.0	1.0

#Removing Outliers

Z-score method

#Pros: Good for normally distributed data.

#Cons: Not suitable for non-normal data; may miss outliers in skewed distributions.

df_encoded_copy2['fnlwgt_zscore'] = stats.zscore(df_encoded_copy2['fnlwgt'])

```
df_encoded_copy2['fnlwgt'] = np.where(df_encoded_copy2['fnlwgt_zscore'].abs() > 3, np.nan, df_encoded_copy2['fnlwgt']) # Replace outliers with
print(df_encoded_copy2.head())
```

```

age  workclass  fnlwgt  education  educational-num  \
0 -0.995125    Private  226802.0      11th              7.0
1 -0.046907    Private   89814.0      HS-grad             9.0
2 -0.776305    Local-gov  336951.0    Assoc-acdm          12.0
3  0.390732    Private  160323.0    Some-college        10.0
4 -1.505704      ?      103497.0    Some-college        10.0

marital-status  occupation  relationship  capital-gain  \
0  Never-married  Machine-op-inspct  Own-child      0
1  Married-civ-spouse  Farming-fishing  Husband      0
2  Married-civ-spouse  Protective-serv  Husband      0
3  Married-civ-spouse  Machine-op-inspct  Husband    7688
4  Never-married      ?              Own-child      0

capital-loss  hours-per-week  native-country  income  gender_Encoded  \
0           0           0.397959  United-States  <=50K      1.0
1           0           0.500000  United-States  <=50K      1.0
2           0           0.397959  United-States  >50K      1.0
3           0           0.397959  United-States  >50K      1.0
4           0           0.295918  United-States  <=50K      0.0

race_AMER-INDIAN-ESKIMO  race_ASIAN-PAC-ISLANDER  race_BLACK  race_OTHER  \
0           0.0           0.0           0.0           0.0
1           0.0           0.0           0.0           0.0
2           0.0           0.0           0.0           0.0
3           0.0           0.0           1.0           0.0
4           0.0           0.0           0.0           0.0

race_WHITE  fnlwgt_zscore
0           0.0           0.419934
1           1.0          -1.017089
2           1.0           1.575412
3           0.0          -0.277440
4           1.0          -0.873553
```

```
#Removing Outliers
```

```
# Median replacement for outliers
```

```
#Pros: Keeps distribution shape intact, useful when capping isn't feasible.
```

```
#Cons: May distort data if outliers represent real phenomena.
```

```
df_encoded_copy3['fnlwgt_zscore'] = stats.zscore(df_encoded_copy3['fnlwgt'])
```