

```
import pandas as pd
import numpy as np
from sklearn import linear_model
import matplotlib.pyplot as plt
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split
```

```
from google.colab import files
uploaded = files.upload()
```

Choose Files salary.csv

- **salary.csv**(text/csv) - 346 bytes, last modified: 3/17/2025 - 100% done

Saving salary.csv to salary.csv

```
df_salary = pd.read_csv('salary.csv')
print(df_salary.head())
```

```
YearsExperience  Salary
0               1.1   39343
1               1.3   46205
2               1.5   37731
3               2.0   43525
4               2.2   39891
```

```
print(df_salary.isnull().sum())
```

```
YearsExperience    2
Salary             0
dtype: int64
```

```
mean_years_experience = df_salary['YearsExperience'].mean()
df_salary['YearsExperience'].fillna(mean_years_experience, inplace=True)
```

<ipython-input-5-6f5a7e93c601>:2: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment. The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values is a copy. For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value, inplace=True)'

```
df_salary['YearsExperience'].fillna(mean_years_experience, inplace=True)
```

```
print(df_salary.isnull().sum())
```

```
YearsExperience    0
Salary             0
dtype: int64
```

```
plt.scatter(df_salary['YearsExperience'], df_salary['Salary'], color='blue')
plt.xlabel('Years of Experience')
plt.ylabel('Salary')
plt.title('Years of Experience vs Salary')
plt.show()
```



```
X_salary = df_salary['YearsExperience'].values.reshape(-1, 1) # Independent variable (Years of Experience)
y_salary = df_salary['Salary'].values # Dependent variable (Salary)
```

```
X_train, X_test, y_train, y_test = train_test_split(X_salary, y_salary, test_size=0.2, random_state=42)
```

```
salary_model = LinearRegression()
salary_model.fit(X_train, y_train)
```



```
LinearRegression
LinearRegression()
```

```
predicted_salary_12_years = salary_model.predict([[12]])
print(f"Predicted salary for 12 years of experience: ${predicted_salary_12_years[0]:.2f}")
```



```
Predicted salary for 12 years of experience: $139500.53
```

```
plt.scatter(df_salary['YearsExperience'], df_salary['Salary'], color='blue')
plt.plot(df_salary['YearsExperience'], salary_model.predict(X_salary), color='red') # Regression line
plt.xlabel('Years of Experience')
plt.ylabel('Salary')
plt.title('Years of Experience vs Salary with Regression Line')
plt.show()
```



